

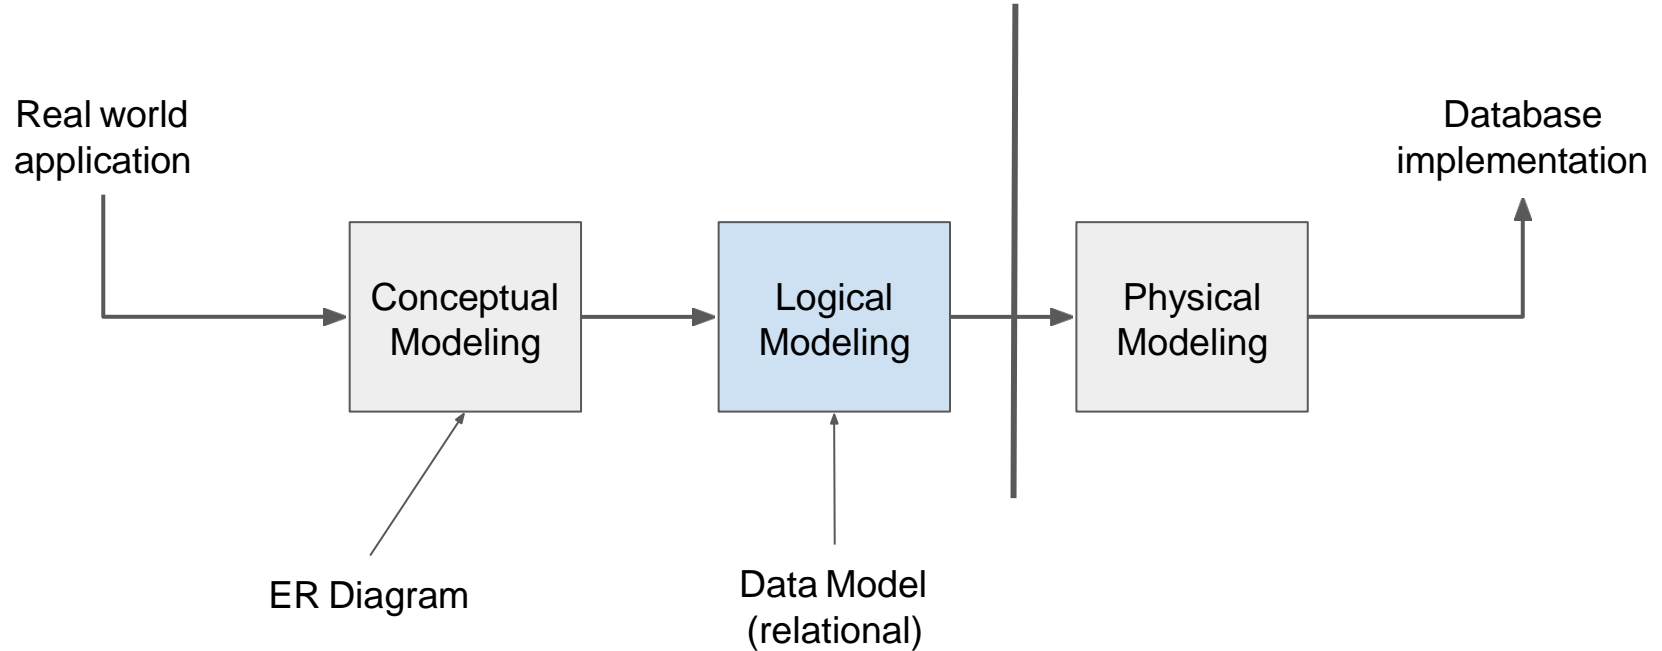
# Backend Development— Complement: NoSQL and MongoDB

Cyrille Jegourel – Singapore University of Technology and Design

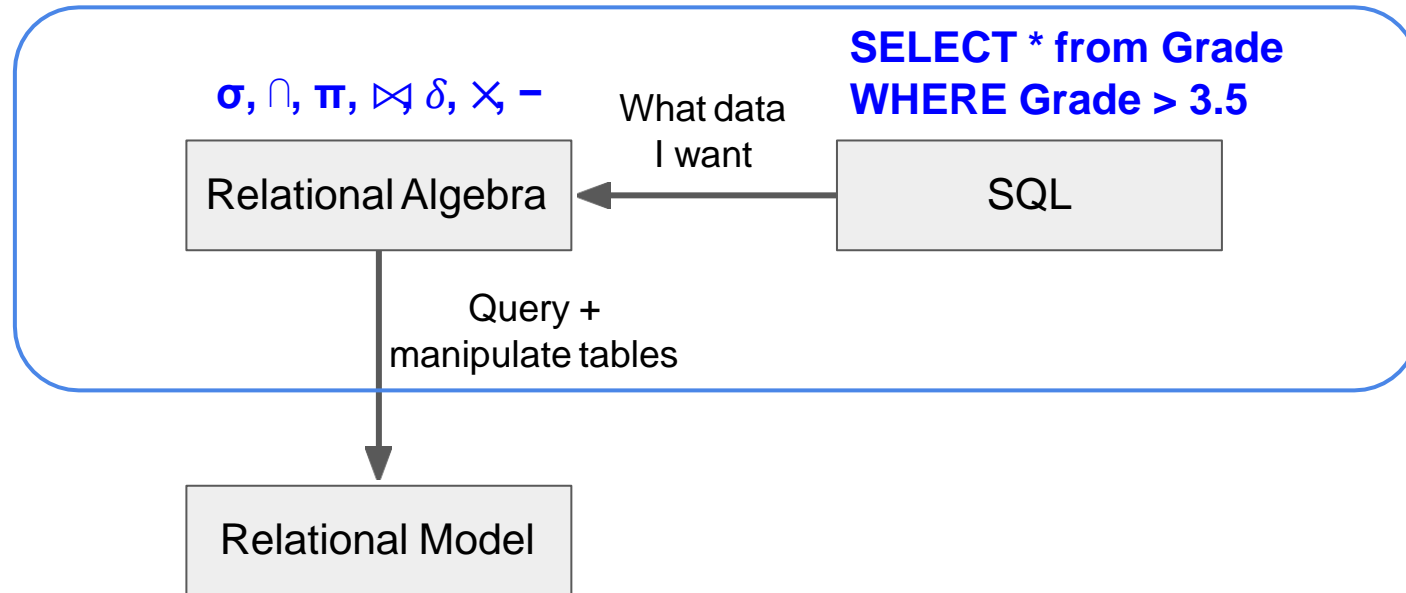


SINGAPORE UNIVERSITY OF  
TECHNOLOGY AND DESIGN

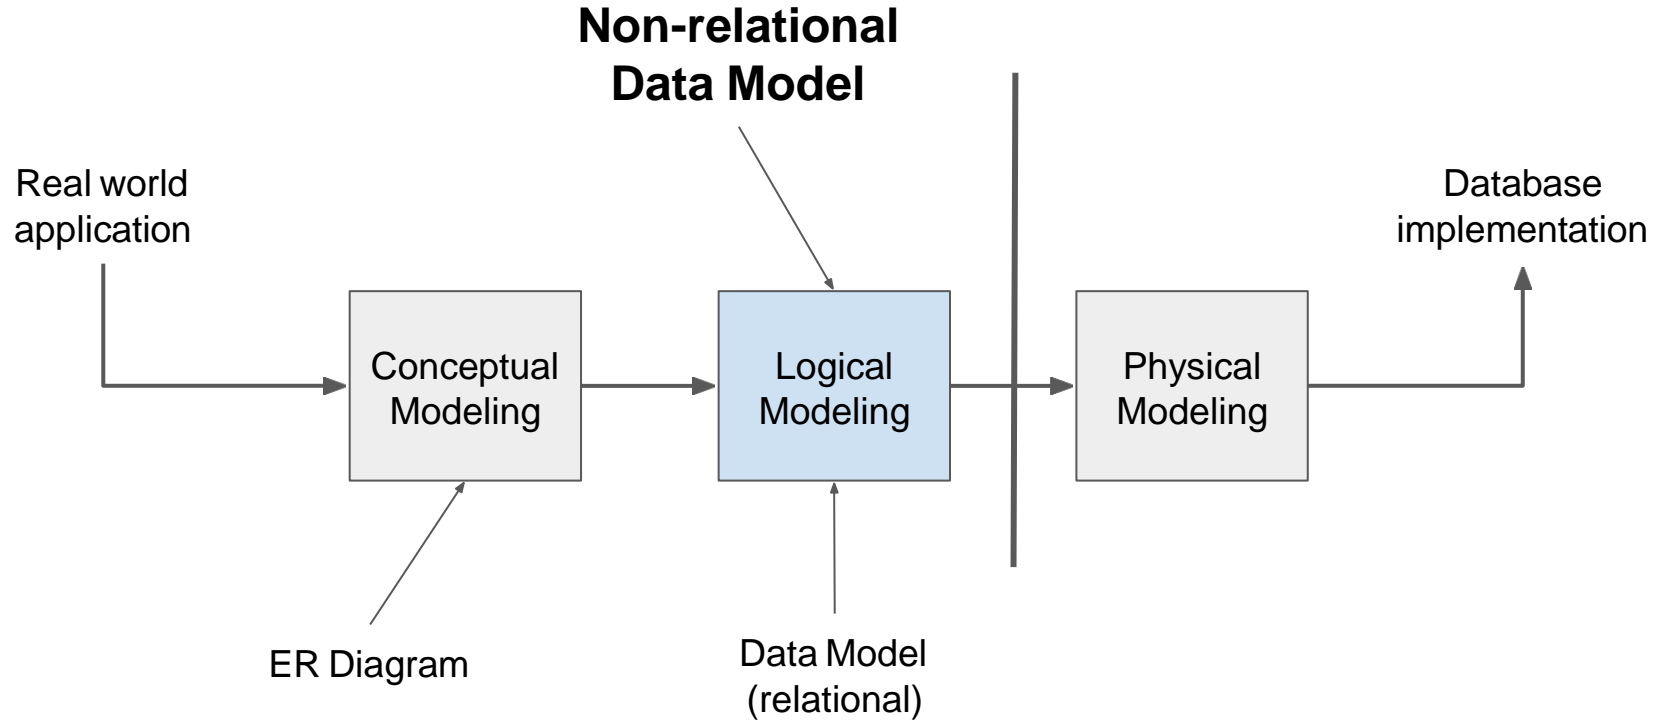
# Recap



# Recap



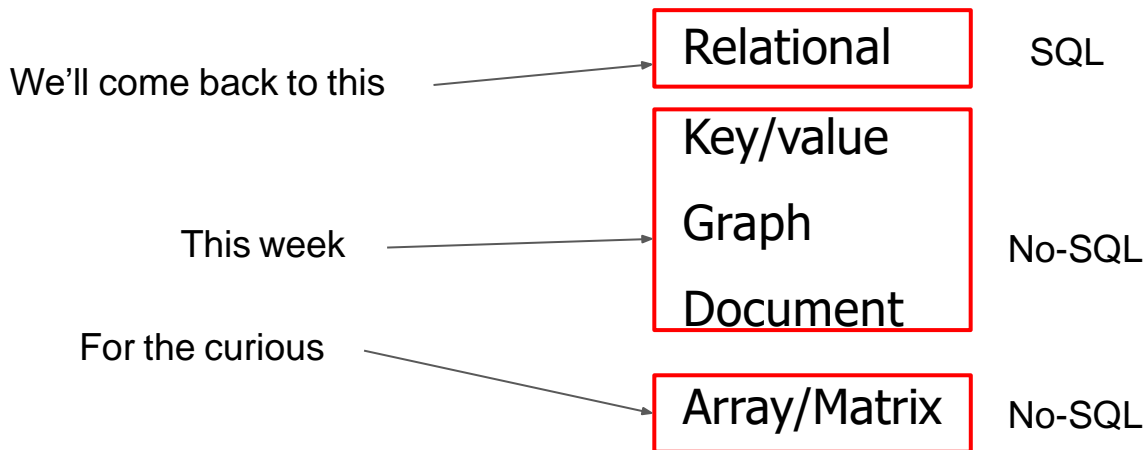
# In these slides



# Data Model - Recap

- Data model:

- How do you **describe** the data to the database?
- In a language the database understand



# Data Model - Recap



Relational

Key/value

Graph

Document

Array/Matrix



NoSQ

# HOW TO WRITE A CV



Leverage the NoSQL boom



# NoSQL

- Started as: no SQL
- Now: Not Only SQL

## Event Details

### Introduction

This meetup is about "open source, distributed, non relational databases".

Have you run into limitations with traditional relational databases? Don't mind trading a query language for scalability? Or perhaps you just like shiny new things to try out? Either way this meetup is for you.

Join us in figuring out why these newfangled Dynamo clones and BigTables have become so popular lately. We have gathered presenters from the most interesting projects around to give us all an introduction to the field.

### Preliminary schedule

09:45: Doors open

10:00: **Intro session** (Todd Lipcon, Cloudera)

10:40: **Voidemort** (Jay Kreps, LinkedIn)

11:20: Short break

11:30: **Cassandra** (Avinash Lakshman, Facebook)

12:10: Free lunch (sponsored by Last.fm)

13:10: **Dynomite** (Cliff Moon, Powerset)

13:50: **HBase** (Ryan Rawson, Stumbleupon)

14:30: Short break

14:40: **Hypertable** (Doug Judd, Zvents)

15:20: **CouchDB** (Chris Anderson, couch.io)

16:00: Short break

16:10: Lightning talks

16:40: Panel discussion

17:00: Relocate to Kate O'Brien's, 579 Howard St. @ 2nd. First round sponsored by Digg

### Registration

The event is free but space is limited, please register if you wish to attend.

### Location

Magma room, CBS interactive

235 Second Street

San Francisco, CA 94105



**thrubd** @thrubd · 23 May 2009

sucks i'm not on the west coast and will not be able to attend #nosql



**Todd Lipcon** @tlipcon · 23 May 2009

working on slides for the #nosql meetup in June. trying to cover all of dist systems in 40 minutes is not as easy as it sounds.



**Chris Anderson** @jchris · 13 May 2009



It's official - I'll be relaxifying everyone with CouchDB at #NoSQL. Thanks @skr!



# Why NoSQL?


- You have a few tables
- Few GBs max


Your data fit on a **single machine**




# Why NoSQL?

- Or a single server




# Why NoSQL?

Petabytes do  
not fit on a  
single  
machine




# Why NOSQL

- Scale up:
  - Bigger machines
  - Cannot count on it forever
    - Moore Law is slowing down (dead!)



# Why NOSQL

- Scale out
  - Many more machines that act as a single machine




# Why NOSQL

- Scale-out is adopted in practice
  - More difficult that it looks

**“You can have a second computer  
once you’ve shown you know how  
to use the first one”**

(Paul Barham)



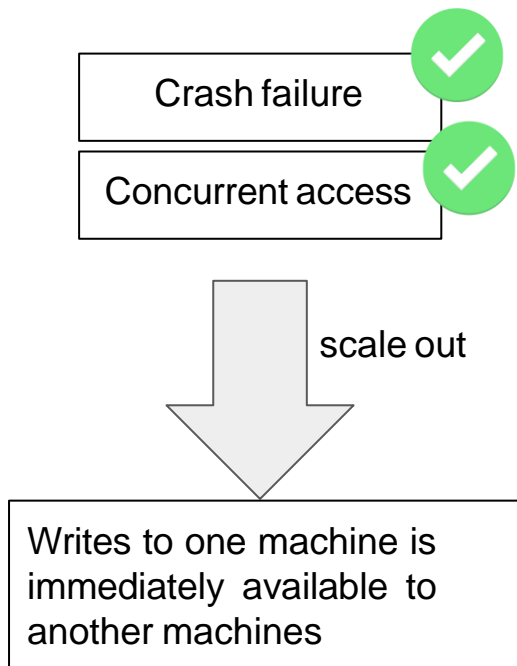
A distributed system is one in which  
the failure of a computer you didn't  
even know existed can render your  
own computer unusable.

— *Leslie Lamport* —

AZ QUOTES

# Why NoSQL

- <Massive hand waving>
  - Very difficult to implement the Relation Model over many machines
  - Check out 50.041 if interested





# Why NoSQL

- Facebook:
  - Don't care if I don't see photos in real-time
  - Care if I can't upload a photo
- Amazon:
  - Don't care if my cart forgot an item
  - Care if I can't add an item



Big companies are  
obsessed about this

# NOSQL




Relational Database

Schema: tables

SQL: join, select, ect.

Correctness

Speed\*



NoSQL Database

No Schema: just blobs

Simple API: put/get

Not always correct


Scalability

Scale out

# Key-Value Store

# Key-Value Store

- Absolute opposite of relational database
- Local:
  - LevelDB
- Remote services
  - Redis
  - Memcached



# Key-Value Store

- Schema

- key  $\rightarrow$  value
- Like a hash table


- Language:

- put(key, value)
- get(key)

K1	AAABBBCCC
K2	12340lks 25/11/2019
K3	.=lkj23ois09320-981
K4	abcdo093kjkap

# LevelDB

- Backend of Chrome!
- What happen:
  - Put(.) on existing key?
  - Get(.) on non-existent keys?
- Demo with ldb tool

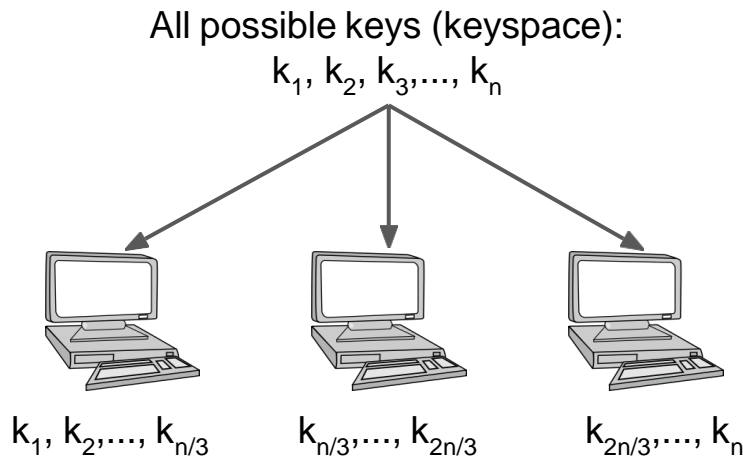


```
%. /ldb db
>ls
apple1  apple2  apple3  apple4  apple5  apple6  apple7  apple8  apple9  foobar1  foobar2
foobar3  foobar4  foobar5  foobar6  foobar7  foobar8  foobar9  quxx    zebra1  zebra2  zebra3
zebra4  zebra5  zebra6  zebra7  zebra8  zebra9
>
zebra1  zebra2  zebra3  zebra4  zebra5  zebra6  zebra7  zebra8  zebra9
>get zebra1
```

# Key-Value Store

- The good

- Very simple to use, behave like a map
- Very handy in many types of applications (Web applications)
- Very fast, like a map
- Very scalable
  - Partition the keys into disjoint sets
  - Then distribute them over many machines



# Key-Value Store

- The bad

- You only have Put(.) and Get(.)
- Range queries, like `select * from Grade where gpa > 3.5;`
- Join queries not supported.


```
for each key in keyspace:  
    if (int)key[value] > 3.5:  
        print(key)
```



# Key-Value Store


- The ugly:
  - Treat everything as a binary string (blob, no structure)
  - If your value contains structure, e.g., (age, name, salary)
    - Too bad for you
  - Push complexity to someone else!
    - Price will be paid, question is where, and by whom
  - Joins and range queries are costly

# Key-Value Store



# More Structured Key-Value Store

- LevelDB:
  - Flat keyspace
  - An object is identified by a single key
  - Like a typical hashtable
- (A bit) more structure
  - Organize keys into groups
  - An object is identified by (group key, object key)
  - Like a 2-level hash table




# S3

- Amazon service
- All data divided in independent buckets
- Each bucket is a key-value store


- `put((bucket_key, object_key), value)`

- `get((bucket_key, object_key))`



# S3

**Loss of 1 object per  
10K years**



## **Durability & Data Protection**

**Q: How durable is Amazon S3?**

Amazon S3 Standard, S3 Standard-IA, S3 One Zone-IA, and S3 Glacier are all designed to provide **99.999999999%** durability of objects over a given year. This durability level corresponds to an average annual expected loss of 0.000000001% of objects. For example, if you store 10,000,000 objects with Amazon S3, you can on average expect to incur a loss of a single object once every 10,000 years. In addition, Amazon S3 Standard, S3 Standard-IA, and S3 Glacier

# RocksDB

- Built on LevelDB
  - Used inside Facebook
- Column families
  - Grouping keys in a column family
- Demo with Rocksdb's Idb tool



# Redis

- In memory vs. disk
  - All data DRAM
  - Very, very fast

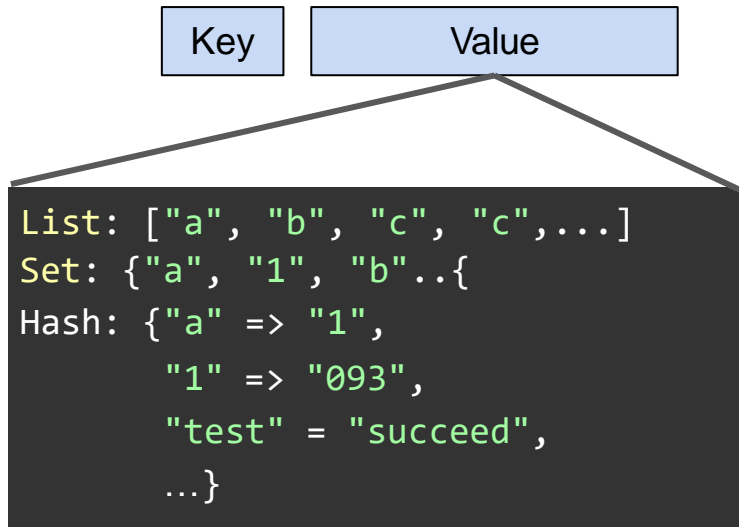


## Numbers Everyone Should Know

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	100 ns
Main memory reference	100 ns
Compress 1K bytes with zippy	10,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from network	10,000,000 ns
Read 1 MB sequentially from disk	30,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

# Redis

- More structure
  - Value has richer type than just blob
  - List
  - Set
  - Even key-value collection





# Redis

- Very popular backend for Web applications
  - Because every ms counts
  - Memcached's cousin
  - Too popular that it's in Ubuntu's standard *apt* repositories



# Redis

- Demo with redis-cli

```
lpush course 12345  
lpush course istd_50043  
llen course  
lindex course 0
```

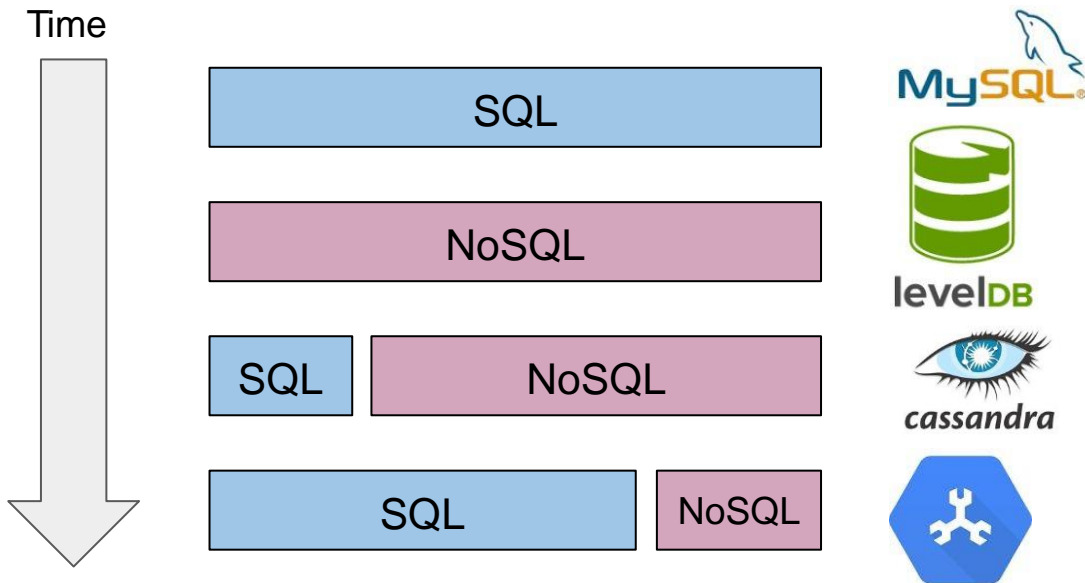
```
hset student age 25  
hset student gpa 3.5  
hvals student  
hkeys student
```

# Summary

- NoSQL chooses scalability over everything else
- Key-value store:
  - LevelDB: bare bone
  - S3/RocksDB: column families
  - Redis: in-memory + remote service + typed values.

# Final Remark

## On History



NoSQL still extremely useful


But people want ease-of-use

- Appreciate the value of SQL more and more

# MongoDB



A database that everybody *loves to hate*



But ends up using it anyway

- Most vocal proponent of the NoSQL movement
  - Among the earliest (2009)
- But the database community weren't impressed



# MongoDB

- Regardless, still the most popular NoSQL



# Document Store

- What is a document?
  - Loose term
  - Anything that is *parse-able*
  - Examples: csv, text, etc.

```
2, San Jose Diridon Caltrain Station,37.3297,-121.902,27, San Jose,2013-08-06,95113
3, San Jose Civic Center,37.3307,-121.889,15, San Jose,2013-08-05,95113
4, Santa Clara at Almaden,37.334,-121.895,11, San Jose,2013-08-06,95113
5, Adobe on Almaden,37.3314,-121.893,19, San Jose,2013-08-05,95113
6, San Pedro Square,37.3367,-121.894,15, San Jose,2013-08-07,95113
7, Paseo de San Antonio,37.3338,-121.887,15, San Jose,2013-08-07,95113
8, San Salvador at 1st,37.3302,-121.886,15, San Jose,2013-08-05,95113
9, Japantown,37.3487,-121.895,15, San Jose,2013-08-05,95113
10, San Jose City Hall,37.3374,-121.887,15, San Jose,2013-08-06,95113
```

```
Mr. Bingley was good-looking and gentlemanlike; he had a pleasant
countenance, and easy, unaffected manners. His sisters were fine women,
with an air of decided fashion. His brother-in-law, Mr. Hurst, merely
looked the gentleman; but his friend Mr. Darcy soon drew the attention
of the room by his fine, tall person, handsome features, noble mien, and
the report which was in general circulation within five minutes
after his entrance, of his having ten thousand a year. The gentlemen
pronounced him to be a fine figure of a man, the ladies declared he
was much handsomer than Mr. Bingley, and he was looked at with great
admiration for about half the evening, till his manners gave a disgust
which turned the tide of his popularity; for he was discovered to be
proud; to be above his company, and above being pleased; and not all
his large estate in Derbyshire could then save him from having a most
forbidding, disagreeable countenance, and being unworthy to be compared
with his friend.
```



# Document Store

- We want semi-structure document:

- Self-explaining documents
- Use tags to capture semantics
- Examples:
  - XML, JSON
  - ProtoBuffer

```
<?xml version="1.0" standalone="yes"?>
<BankAccount>
  <Number>1234</Number>
  <Type>Checking</Type>
  <OpenDate>11/04/1974</OpenDate>
  <Balance>25382.20</Balance>
  <AccountHolder>
    <LastName>Singh</LastName>
    <FirstName>Darshan</FirstName>
  </AccountHolder>
</BankAccount>
```



Tell **us** what the enclosed data means

# Document Store

- Semi-structured documents




```
message Block {  
    uint32 version = 1;  
    google.protobuf.Timestamp timestamp = 2;  
    repeated Transaction transactions = 3;  
    bytes stateHash = 4;  
    bytes previousBlockHash = 5;  
    bytes consensusMetadata = 6;  
    NonHashData nonHashData = 7;  
}  
  
// Contains information about the blockchain le  
// block hash, and previous block hash.  
message BlockchainInfo {  
  
    uint64 height = 1;  
    bytes currentBlockHash = 2;  
    bytes previousBlockHash = 3;  
}
```

# Document Store

- Semi-structured documents

## JavaScript Object Notation (JSON)

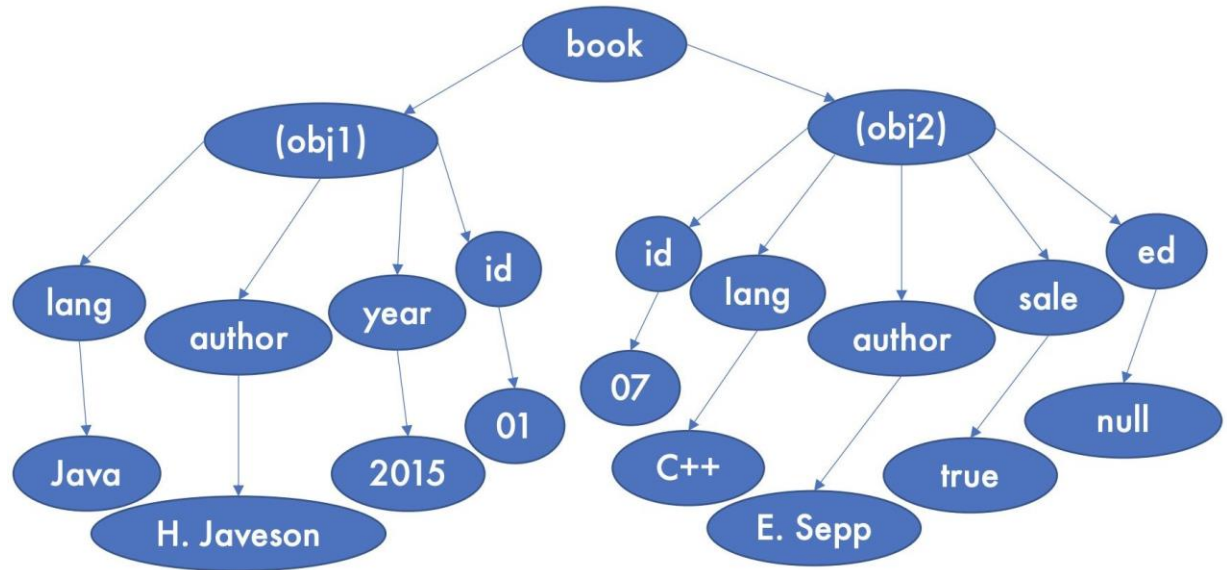


Many applications phasing out  
XML in favor of JSON

```
{  
  "orders": [  
    {  
      "orderno": "748745375",  
      "date": "June 30, 2088 1:54:23 AM",  
      "trackingno": "TN0039291",  
      "custid": "11045",  
      "customer": [  
        {  
          "custid": "11045",  
          "fname": "Sue",  
          "lname": "Hatfield",  
          "address": "1409 Silver Street",  
          "city": "Ashland",  
          "state": "NE",  
          "zip": "68003"  
        }  
      ]  
    }  
  ]  
}
```

# Semi-Structured Documents


- It's a **tree**



# Semi-Structured Documents

- From tables to trees
  - Not hard, because tables = flat trees


Name	Phone
Anh	12345
Dan	23093
Leo	09470



```
{
  "person": [
    {
      "name": "Anh",
      "phone": "12345"
    },
    {
      "name": "Dan",
      "phone": "23093"
    },
    {
      "name": "Leo",
      "phone": "09470"
    }
  ]
}
```

# Semi-Structured Documents

Name	Phone
Anh	12345
Dan	23093
Leo	09470
Ben	NULL




```
{
  "person": [
    {
      "name": "Anh",
      "phone": "12345"
    },
    {
      "name": "Dan",
      "phone": "23093"
    },
    {
      "name": "Leo",
      "phone": "09470"
    },
    {
      "name": "Ben"
    }
  ]
}
```

**Field missing  
= NULL**

# Semi-Structured Documents

- But can we fit any tree into a table?
- Non-flat data:
  - **Array**
  - Multi-part

Name	Phone
Anh	???
Dan	23093
Leo	09470




```
{
  "person": [
    {
      "name": "Anh",
    },
    {
      "name": "Dan",
      "phone": "23093"
    },
    {
      "name": "Leo",
      "phone": "09470"
    }
  ]
}
```

# Semi-Structured Documents

- But can we fit any tree into a table?
- Non-flat data:
  - Array
  - **Multi-part**

Name	Phone
???	12345
Dan	23093
Leo	09470



```
{
  "person": [
    {
      "phone": "12345"
    },
    {
      "name": "Dan",
      "phone": "23093"
    },
    {
      "name": "Leo",
      "phone": "09470"
    }
  ]
}
```



# Semi-Structured Documents

- Relational data model isn't designed for nested data
  - Tables vs. trees
- Term: **impedance mismatch**



# Document Store

- Handle trees
- Many implementations



# MongoDB Language

- Remember for SQL, we have:
  - Data Definition Language (DDL):  
create, delete, etc.
  - Data Manipulation Language (DML):  
insert, update, etc.
  - Query Language: select ... from ...  
where
- MongoDB supports same categories
  - **Tables** in SQL → **Collections** in MongoDB

MySQL	MongoDB
Database	Database
Tables	Collections

```
{
  "Name": "Anh"
  "Phone": [
    "12345",
    "93932"
  ]
}

{
  "Name": "Dan"
  "Phone": "93752"
}
```

Collection

```
{
  "Name": {
    "First": "Albert",
    "Last": "Einstein"
  },
  "Theory": "Particle Physics"
}

{
  "Name": {
    "First": "Sheldon",
    "Last": "Copper"
  },
}
```

Document

Database

# MongoDB

- Create table

```
use university;
```

db (default database)

db ← university

- Create collections

```
db.createCollection("faculty")  
db.createCollection("student")
```

collection name

# MongoDB

- Insert new document
  - Duplicates are allowed

```
db.faculty.insert({"Name": "Einstein",  
"Theory": "Relativity"})
```

```
db.faculty.update({"Name": "Einstein"},  
{"$set": {"Country": "Germany"}})
```

```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "Theory": "Particle Physics"  
}  
  
{  
  "Name": "Einstein",  
  "Theory": "Relativity"  
}  
  
...  
...
```

Guess what this does

# MongoDB

- Read
  - Select all documents

```
db.faculty.find({})
```

```
db.faculty.find()
```

```
SELECT * from University;
```

(SQL CheatSheet)

# MongoDB

- Read condition

```
db.faculty.find(  
  {"Theory": "Particle Physics"}  
)
```

```
SELECT * from Faculty  
WHERE Theory = "Particle Physics"
```

(SQL CheatSheet)

```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "Theory": "Particle Physics"  
}
```



```
{  
  "Name": {  
    "First": "Kurt",  
    "Last": "Godel"  
  },  
  "Theory": "Incompleteness"  
}
```



```
{  
  "Name": {  
    "First": "Sheldon",  
    "Last": "Copper"  
  },  
}
```





# MongoDB

- Read projection

```
db.faculty.find(  
  {"Theory": "Particle Physics"},  
  {"Name": 1, "Last": 1}  
)
```

SELECT Name, Last from Faculty  
WHERE Theory = "Particle Physics"

(SQL CheatSheet)

```
{  
  "First": "Albert",  
  "Last": "Einstein",  
  "Theory": "Particle Physics"  
}  
{  
  "Name": "Higgs",  
  "Theory": "Particle Physics"  
}  
{  
  "First": "Kurt",  
  "Last": "Godel",  
  "Theory": "Incompleteness"  
}
```




```
{ "Last": "Einstein" }  
{ "Name": "Higgs" }
```

# MongoDB

- Read
  - Nested field

```
db.faculty.find({"Name.First": "Albert"})
```






```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "Theory": "Particle Physics"  
}  
  
{  
  "Name": {  
    "First": "Kurt",  
    "Last": "Godel"  
  },  
  "Theory": "Incompleteness"  
}  
  
{  
  "Name": {  
    "First": "Sheldon",  
    "Last": "Copper"  
  },  
}
```

# MongoDB

- Read
  - Nested field

```
db.faculty.find({"Name": {"First": "Albert"}})
```

```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "Theory": "Particle Physics"  
}  
  
{  
  "Name": {  
    "First": "Kurt",  
    "Last": "Godel"  
  },  
  "Theory": "Incompleteness"  
}  
  
{  
  "Name": {  
    "First": "Albert",  
  },  
  "Theory": "Unification"  
}
```



# MongoDB

- Read
  - List matching

```
db.faculty.find({"Theory": "Special  
relativity"})
```

```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "Theory": [  
    "Special relativity",  
    "General relativity"  
  ]  
}  
  
{  
  "Name": "Godel",  
  "Theory": "Incompleteness"  
}
```



# MongoDB

- Read operators

Operator	Meaning
\$gte, \$eq, \$ne, \$gt, \$lt, \$lte	>=, =, !=, >, <, <=
\$in, \$nin	∈, ∉

```
db.faculty.find( {  
  "NoPublications": {"$gte": 120}  
})
```

```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "NoPublications": 209  
}  
  
{  
  "Name": "Godel",  
  "NoPublications": 100  
}
```



# MongoDB

- Read operators

Operator	Meaning
\$gte, \$eq, \$ne, \$gt, \$lt, \$lte	>=, =, !=, >, <, <=
\$in, \$nin	∈, ∉

```
db.faculty.find( {  
  "University": {"$in": ["NUS", "SUTD"]}  
})
```

```
{  
  "Name": {  
    "First": "Albert",  
    "Last": "Einstein"  
  },  
  "NoPublications": 209  
  "University": "SUTD"  
},  
{  
  "Name": "Godel",  
  "NoPublications": 100  
  "University": "SMU"  
}
```



# MongoDB

- Read

- Count

```
db.faculty.find({"Theory": "Particle Physics"}).count()
```

- Sort

```
db.faculty.find({"Theory": "Particle Physics"}).sort({"Age": -1})
```

- Limit

- Duplicates

```
db.faculty.distinct("Name", {"Theory": "Particle Physics"})
```

- Check them out yourself

# Summary

- MongoDB most popular for
  - Semi-structured data model
  - Or trees
- One size doesn't fit all

## “One Size Fits All”: An Idea Whose Time Has Come and Gone

Michael Stonebraker  
Computer Science and Artificial  
Intelligence Laboratory, M.I.T., and  
StreamBase Systems, Inc.  
stonebraker@csail.mit.edu

Uğur Çetintemel  
Department of Computer Science  
Brown University, and  
StreamBase Systems, Inc.  
ugur@cs.brown.edu

### Abstract

*The last 25 years of commercial DBMS development can be summed up in a single phrase: “One size fits all”. This phrase refers to the fact that the traditional DBMS architecture (originally designed and optimized for business data processing) has been used to support many data-centric applications with widely varying characteristics and requirements.*

of multiple code lines causes various practical problems, including:

- *a cost problem*, because maintenance costs increase at least linearly with the number of code lines;
- *a compatibility problem*, because all applications have to run against every code line;
- *a sales problem*, because salespeople get confused about which product to try to sell to a customer; and
- *a marketing problem*, because multiple code lines